

Predicting Cross-Sectional Stock Returns: A Machine Learning Approach using Kernel Ridge Regression and Mean-Variance Optimization

Antonio Gabriele Santini

Master's Degree in Quantitative Finance UniTo

July 2026

Abstract

This paper investigates the predictability of cross-sectional stock returns using Kernel Ridge Regression (KRR) and evaluates its economic value through a Mean-Variance portfolio allocation framework. We construct a deep historical panel dataset comprising ten major US equities from 1990 to 2024, calculating eight price- and volume-based characteristics that are cross-sectionally standardized to prevent look-ahead bias. To capture asset-specific behaviors and mitigate overfitting, we train distinct, independent KRR models for each asset and systematically evaluate Linear, Polynomial, and Gaussian (RBF) kernels. We implement a rigorous Walk-Forward (Rolling Window) backtesting methodology, coupled with dynamic Bayesian hyperparameter optimization applied iteratively over a 21-year out-of-sample period (January 2003 to November 2024). Crucially, to bridge the gap between theoretical forecasts and realistic operational implementation, we evaluate all algorithmic portfolios net of a 10 basis points transaction cost. Our results reveal a stark bias-variance tradeoff compounded by market frictions: the rigid Linear Kernel proved to be the most robust predictive model (Gross Return 1628.67%, Net Return 1396.80%, Annualized Net Sharpe 1.0121), while highly parameterized non-linear models severely underperformed due to noise-driven overfitting and excessive portfolio turnover (up to 69.40% monthly). Ultimately, the passive, zero-turnover Equal-Weight (1/N) benchmark dominated all active ML strategies with an Annualized Net Sharpe Ratio of 1.2055, highlighting the profound impact of survivorship bias and the friction drag of active allocation in highly efficient mega-cap equity markets.

Keywords: Machine Learning, Kernel Ridge Regression, Asset Pricing, Walk-Forward Optimization, Transaction Costs, Survivorship Bias.

1. INTRODUCTION

Financial markets are highly complex, non-stationary systems characterized by a notoriously low signal-to-noise ratio. Forecasting cross-sectional stock returns remains one of the most formidable challenges in quantitative finance. Traditional linear asset pricing frameworks often operate under restrictive functional form assumptions, struggling to capture the intricate interactions among firm characteristics that drive future returns. Consequently, the quantitative investment industry has increasingly turned to Machine Learning (ML) techniques, which offer the flexibility to dynamically extract complex predictive patterns directly from the data.

The primary objective of this paper is to rigorously test whether advanced ML methodologies—specifically Kernel Ridge Regression (KRR)—can overcome market noise and deliver superior out-of-sample economic value compared to a naive asset allocation framework. To this end, we construct regularized mean-variance portfolios based on the expected return predictions generated by Linear, Polynomial, and Gaussian Radial Basis Function (RBF) KRR models. We evaluate these dynamically weighted ML portfolios against the robust "Equal-Weight" ($1/N$) strategy, a benchmark widely recognized in the financial literature for its empirical resilience and ability to outperform many sophisticated optimization techniques.

To conduct this analysis, we process a deep historical panel dataset of ten major U.S. equities spanning over three decades (1990–2024), computing eight distinct price- and volume-based predictive characteristics. A time-series-aware methodology is strictly enforced throughout the study: we implement cross-sectional standardization to eliminate look-ahead bias and estimate independent, asset-specific KRR models to capture idiosyncratic behaviors. Crucially, we abandon static Train/Test splits in favor of a robust Walk-Forward (Rolling Window) optimization over a 21-year out-of-sample period (2003–2024). Within this moving window, we employ Bayesian hyperparameter tuning via the Tree-structured Parzen Estimator (TPE) algorithm, dynamically recalibrating the models to shifting market regimes.

A defining feature of this study is the integration of realistic market frictions. Academic literature frequently reports algorithmic outperformance based strictly on gross returns, ignoring the severe performance drag induced by portfolio turnover. In this paper, all active ML portfolios are evaluated net of a realistic 10 basis points transaction cost, allowing for an uncompromising assessment of their true operational viability.

The empirical findings from this extended out-of-sample framework provide a compelling lesson on the bias-variance tradeoff in financial machine learning. Contrary to expectations derived from static backtests, the simpler Linear Kernel emerged as the most resilient predictive model, outperforming the highly flexible Polynomial and Gaussian kernels. Over a 10-year rolling window, highly parameterized non-linear models tend to overfit the idiosyncratic noise of a concentrated 10-asset universe. This overfitting manifests not only in poor predictive accuracy but in erratic allocations, leading to massive portfolio turnover (averaging 69.40% monthly for the Polynomial model) that rapidly destroys alpha through transaction costs.

However, underscoring the enduring difficulty of active portfolio allocation, the simple Equal-Weight benchmark eclipsed all active ML portfolios in risk-adjusted terms, yielding the highest overall Annualized Net Sharpe Ratio of 1.2055 and a Net Cumulative Return of 2679.07%. These results suggest that while ML can extract predictive signals, overcoming the zero-turnover efficiency, the parameter-free diversification, and the structural survivorship bias of a passive $1/N$ long-only strategy remains a monumental challenge in quantitative asset management.

2. DATA AND FEATURE ENGINEERING

2.1. Asset Universe and Historical Context

The empirical analysis is conducted on a deliberately selected universe of ten highly liquid, large-capitalization U.S. equities: Microsoft (MSFT), JPMorgan Chase (JPM), Johnson & Johnson (JNJ), Procter & Gamble (PG), McDonald's (MCD), ExxonMobil (XOM), Caterpillar (CAT), Apple (AAPL), AT&T (T), and Walmart (WMT).

We extract monthly adjusted closing prices and trading volumes from January 1990 to November 2024. The primary justification for utilizing such deep historical data is to train and evaluate the Machine Learning algorithms across a wide array of distinctly different macroeconomic regimes. A backtest spanning over three decades successfully captures major structural breaks and volatility clusters, including the Dot-Com Bubble (2000-2002), the Global Financial Crisis (2008-2009), and the COVID-19 pandemic (2020).

From a quantitative implementation perspective, this universe serves two purposes. First, the extreme liquidity of these mega-cap equities ensures minimal bid-ask spreads and market impact, structurally justifying the application of a rigorous but manageable 10 basis points transaction cost during the backtesting phase. Second, while this static selection ensures computational tractability, it intentionally exposes the framework to *survivorship bias*—testing exclusively on companies that survived and thrived over a 34-year horizon. Acknowledging this bias is crucial for a realistic evaluation of active vs. passive (Equal-Weight) asset allocation dynamics.




Figure 1: Raw Data: Asset Growth (Base 100 Normalized). The chart illustrates the cumulative price evolution of the 10 selected assets from 1990 to 2024, highlighting the exponential growth of technology components (e.g., AAPL, MSFT) and the divergent volatility profiles across the universe.

2.2. Feature Construction

To map the inputs for the Kernel Ridge Regression, we engineer a comprehensive set of eight predictive characteristics ($c_1, \dots, c_8$) strictly derived from historical prices and trading volumes. By explicitly excluding fundamental accounting data, we entirely bypass the latency and revision issues typically associated with the publication of corporate balance sheets, ensuring pure out-of-sample tractability. For each asset i at month t , we compute the following quantitative factors:

- **Short-Term Reversal (C_1):** The simple 1-month return $R_{i,t}$, capturing the tendency of assets to mean-revert over short horizons due to temporary liquidity shocks or micro-structural overreactions.
- **12-Month Momentum (C_2):** The cumulative return from month $t - 12$ to $t - 1$, computed as $(P_{i,t-1}/P_{i,t-12}) - 1$. This captures the classic cross-sectional trend-following anomaly.
- **6-Month Momentum (C_3):** A medium-term trend metric, computed as the cumulative return from month $t - 6$ to $t - 1$, capturing intermediate price persistence.
- **Historical Volatility (C_4):** The rolling 12-month standard deviation of monthly returns, $\sigma_{i,(t-12 \rightarrow t)}$. In empirical asset pricing, this acts as a primary proxy for the asset's idiosyncratic risk.
- **Log Volume (C_5):** The natural logarithm of the monthly trading volume, $\ln(V_{i,t})$, serving as a widely accepted proxy for investor attention and absolute liquidity.
- **Amihud Illiquidity Proxy (C_6):** Calculated as the ratio of absolute monthly return to trading volume, $|R_{i,t}|/V_{i,t}$. This measures price impact, specifically identifying how much the asset's price moves per unit of traded volume.
- **Long-Term Reversal (C_7):** The cumulative return evaluated from month $t - 36$ to $t - 13$. Empirical literature demonstrates that extreme multi-year winners often face structural corrections, becoming subsequent losers.
- **Max Return / Lottery Effect (C_8):** The maximum single-month return achieved over the trailing 12 months. This captures the behavioral finance anomaly where retail investors overpay for "lottery-like" stocks with right-skewed return distributions.

2.3. Target Variable and Bias Prevention via Cross-Sectional Standardization

The predictive target for the supervised machine learning algorithms is the forward-looking 1-month return, defined as $R_{i,t+1}$.

A critical methodological imperative in algorithmic backtesting is the absolute prevention of *look-ahead bias* (data leakage). If a model scales a characteristic using the global mean and variance across the entire 1990-2024 dataset, it inadvertently injects future information into the historical training phase.

To structurally prohibit this, we enforce strict Cross-Sectional Standardization. At each specific time step t , the characteristic c for asset i is standardized exclusively against the cross-sectional mean ($\mu_{c,t}$) and standard deviation ($\sigma_{c,t}$) of the available universe at that exact historical moment:

$$Z_{i,c,t} = \frac{X_{i,c,t} - \mu_{c,t}}{\sigma_{c,t} + \epsilon} \quad (1)$$

where ϵ is a minimal constant to prevent division by zero. This operation converts the raw features into Z-scores ($Z_{i,t}$), ensuring that the machine learning algorithm solely evaluates the *relative* cross-sectional positioning of an asset (e.g., "how high is AAPL's momentum compared to the rest of the universe today?") using only the information set strictly available at time t .

Furthermore, this standardization is computationally essential for Kernel Ridge Regression. Because kernel functions (particularly the Gaussian RBF) rely on Euclidean distance metrics, scaling ensures that features with naturally large absolute values (such as unscaled Volume) do not disproportionately dominate the kernel similarity matrix computations.

3. METHODOLOGY

3.1. Kernel Ridge Regression (KRR)

To model the highly complex relationship between firm characteristics and future returns, we employ Kernel Ridge Regression (KRR). Traditional linear predictive regressions frequently suffer from underfitting when confronted with the non-linear realities of financial markets. KRR bypasses this limitation by utilizing the "kernel trick," an elegant mathematical mechanism that implicitly maps the original characteristic space into a high-dimensional (potentially infinite) feature space where the data becomes linearly separable, without the computational burden of explicitly computing the transformations.

To account for the idiosyncratic behavior of each asset and mitigate the statistical noise associated with heterogeneous panel data, **we train a distinct and independent KRR model for each individual asset i .** Let $Z_{train,i}$ be the $T \times K$ matrix of cross-sectionally standardized characteristics observed over the training window for asset i , and $R_{train,i}$ be the $T \times 1$ vector of its corresponding realized returns. In the dual formulation of Ridge Regression for a specific asset, the objective is to estimate an asset-specific parameter vector α_i rather than a traditional global feature weight vector β . The optimal solution is given by:

$$\hat{\alpha}_i = (K_i + \lambda I_T)^{-1} R_{train,i} \quad (2)$$

where K_i is the $T \times T$ Kernel similarity matrix strictly for asset i , I_T is the identity matrix, and $\lambda > 0$ is the Ridge penalty parameter (L2 regularization) that shrinks coefficients to prevent overfitting on market noise.

Once the individual models are trained, the out-of-sample expected returns for the forward month ($\hat{R}_{test,i}$) are predicted by computing the kernel mapping between the new out-of-sample observations ($Z_{test,i}$) and the historical training data ($Z_{train,i}$):

$$\hat{R}_{test,i} = K_{test,i}(Z_{test,i}, Z_{train,i}) \hat{\alpha}_i \quad (3)$$

To rigorously test whether non-linearities add economic value out-of-sample, we implement and compare three distinct kernel functions:

- **Linear Kernel (Baseline):** $k(z, z') = z^\top z'$

This represents the standard inner product, making the KRR mathematically equivalent to a standard Linear Ridge Regression. It serves as our mandatory baseline to isolate the pure value of non-linear mappings.

- **Polynomial Kernel:** $k(z, z') = (c + z^\top z')^d$

This kernel explicitly models interactions between different characteristics up to degree d . The hyperparameter c controls the trade-off between higher-order and lower-order interaction terms.

- **Gaussian Radial Basis Function (RBF) Kernel:** $k(z, z') = \exp(-c\|z - z'\|^2)$

A highly flexible, non-linear kernel capable of mapping inputs into an infinite-dimensional space. The hyperparameter c dictates the bandwidth of the Gaussian curve, determining the radius of influence of the historical training points.

3.2. Portfolio Construction and Optimization

Generating accurate return forecasts is only the first step in quantitative asset management; the forecasts must be operationalized into a tradable portfolio. To achieve this, we utilize a regularized

Mean-Variance optimization framework dynamically applied at each rolling window step.

Using the historical data within the current training window, we estimate the sample covariance matrix $\hat{\Sigma}$. Given the inherent estimation error in financial covariance matrices, computing the direct inverse $\hat{\Sigma}^{-1}$ often leads to extreme, unstable portfolio weights. To ensure numerical stability and invertibility, we apply a Ridge regularization penalty γ to the diagonal of the covariance matrix. The raw optimal weights w_t^{raw} for the forward month t are computed as:

$$w_t^{raw} = (\hat{\Sigma} + \gamma I_N)^{-1} \hat{R}_t \quad (4)$$

where I_N is the $N \times N$ identity matrix and $\hat{R}_t$ is the vector of the combined asset-specific machine learning return forecasts.

3.2.1 Gross Exposure Normalization

A fundamental issue arises when evaluating portfolios optimized with dynamic regularization parameters: different values of γ inevitably scale the absolute magnitude of the raw weights, leading to drastic variations in the portfolio's gross exposure (leverage). A model with a high γ might invest only a fraction of the available capital, artificially deflating its volatility and cumulative return when compared to a fully invested benchmark.

To ensure a methodologically sound comparison with the Equal-Weight benchmark (which systematically utilizes 100% of the capital), we impose a strict Gross Exposure Normalization constraint. We force the sum of the absolute weights (both long and short positions) to equal exactly 1 (100% capital utilization). The final tradable weights w_t are obtained via:

$$w_t = \frac{w_t^{raw}}{\sum_{i=1}^N |w_{i,t}^{raw}|} \quad (5)$$

By implementing this constraint, we guarantee that the performance differentials between the ML models and the $1/N$ strategy are purely the result of superior active stock selection, rather than arbitrary differences in leverage.

3.2.2 Turnover and Transaction Costs

In realistic algorithmic trading environments, changing portfolio allocations incurs market frictions (brokerage fees, bid-ask spread crossing, and market impact). Ignoring these costs frequently leads to the overestimation of highly active strategies. We explicitly model portfolio turnover as the sum of absolute changes in target weights from month $t - 1$ to month t :

$$Turnover_t = \sum_{i=1}^N |w_{i,t} - w_{i,t-1}| \quad (6)$$

To evaluate the true operational viability of the machine learning forecasts, we apply a stringent transaction cost (c) of 10 basis points (0.0010) per unit of turnover. The realized Net Return ($R_{net,t}$) of the portfolio is thus computed as:

$$R_{net,t} = w_t^\top R_{realized,t} - (Turnover_t \times c) \quad (7)$$

4. HYPERPARAMETER TUNING AND MODEL VALIDATION

4.1. Walk-Forward (Rolling Window) Methodology

To robustly evaluate our models and reflect the real-world operational constraints of quantitative asset management, we abandon the static Train-Validation-Test split in favor of a rigorous Walk-Forward (Rolling Window) optimization. Financial markets are fundamentally non-stationary and subject to frequent structural regime shifts; a static model trained solely on data from the 1990s would inevitably suffer from severe performance decay when confronted with modern market dynamics.

Our implementation dynamically re-trains the asset-specific KRR models at each time step t . Specifically, we employ a rolling window of 120 months (10 years) to construct the historical similarity matrix K_i and the sample covariance matrix $\hat{\Sigma}$. The models generate out-of-sample expected returns $\hat{R}_{t+1}$ exclusively for the single forward month $t + 1$. The window is then advanced by one month, incorporating the newest realized data and dropping the oldest observation. This dynamic out-of-sample design completely eradicates data leakage, respects the chronological autocorrelation of financial time-series, and allows the models to continuously adapt to evolving macroeconomic regimes.

4.2. Bayesian Optimization via TPE Algorithm

Model selection and hyperparameter tuning are computationally intensive tasks, particularly when evaluating Kernel Ridge Regressions across hundreds of rolling windows. A traditional exhaustive Grid Search suffers heavily from the curse of dimensionality and allocates vast computational resources to evaluating evidently sub-optimal parameter sub-spaces.

To maximize both computational efficiency and out-of-sample accuracy, we deploy a Bayesian Optimization framework utilizing the Tree-structured Parzen Estimator (TPE) algorithm, implemented via the *Optuna* library. To prevent extreme computational overhead during the walk-forward backtest, the Bayesian optimizer is triggered periodically (every 12 months). During this recalibration phase, the most recent 24 months within the current 120-month rolling window are isolated as a validation set.

Crucially, to align the machine learning objective with our operational constraints, the TPE algorithm constructs a probabilistic surrogate model explicitly designed to maximize the **Annualized Net Sharpe Ratio** on this validation subset. By incorporating the 10 basis points transaction cost directly into the optimization objective, the algorithm organically learns to penalize hyperparameter configurations that generate excessive turnover, dynamically favoring higher covariance regularization (γ) when trading friction outweighs predictive alpha.

Over 20 intelligent trials per calibration, we optimize the following dimensions simultaneously:

- **λ (Ridge Penalty):** Explored on a continuous logarithmic scale ($\in [0.01, 100.0]$) to control KRR overfitting on asset-specific noise.
- **γ (Covariance Regularization):** Searched logarithmically ($\in [0.01, 50.0]$) to ensure well-conditioned portfolio weights and actively suppress high-turnover allocations.
- **c and d (Kernel Parameters):** Dynamically tuned based on the selected kernel to capture optimal feature interactions and bandwidths.

To adhere to strict academic standards, a dynamic deterministic seed is enforced within the TPE sampler for each window, ensuring that all empirical results presented in this paper remain 100% reproducible.

5. OUT-OF-SAMPLE RESULTS AND DISCUSSION

5.1. Predictive Performance and the Friction of Transaction Costs

The definitive test of any quantitative investment strategy is its out-of-sample empirical performance, net of market execution costs. By utilizing a 120-month rolling window, our out-of-sample testing period spans a rigorous 21-year horizon (January 2003 to November 2024). This period encompasses multiple severe volatility clusters and structural regime shifts, including the 2008 Global Financial Crisis, the European Sovereign Debt Crisis, and the 2020 COVID-19 pandemic shock.

To bridge the gap between theoretical predictive accuracy and operational reality, we penalize all algorithmic rebalancing with a conservative but realistic transaction cost of 10 basis points (0.10%) per unit of turnover. Table 1 presents the core evaluation metrics for the dynamically allocated ML portfolios compared to the naive diversification strategy.

Table 1: Out-of-Sample Performance Summary (2003-01 to 2024-11, Net of 10 bps Fees)

Model	Gross Return	Net Return	Ann. Sharpe	Annual Volatility	Mean Turnover
Linear Kernel	1628.67%	1396.80%	1.0121	13.10%	55.27% / mo
Polynomial Kernel	723.60%	586.72%	0.6987	14.04%	69.40% / mo
Gaussian RBF	1055.07%	903.38%	0.8661	13.21%	53.93% / mo
Equal-Weight (1/N)	-	2679.07%	1.2055	13.41%	0.00% / mo

The trajectory of the resulting wealth generation is visually captured in Figure 2, illustrating the comparative equity curves evaluated strictly net of trading frictions over the 21-year period.




Figure 2: Out-of-Sample Cumulative Returns (Net of Transaction Costs). The graph compares the net wealth generation of the dynamically optimized KRR models against the Equal-Weight benchmark from 2003 to 2024. The massive outperformance of the passive benchmark highlights the combined impact of zero-turnover efficiency and survivorship bias in a long-only mega-cap universe.

5.2. The Bias-Variance Tradeoff and Portfolio Turnover Drag

Analyzing the empirical results in Table 1, the Walk-Forward methodology exposes a classic *bias-variance tradeoff*, which in financial applications manifests directly as portfolio turnover drag.

The **Linear Kernel** emerged as the most robust active predictive model, achieving a Gross Cumulative Return of 1628.67%. However, extracting this alpha required a mean monthly turnover of 55.27%, meaning the algorithm aggressively rotated more than half of its capital every 30 days. This relentless rebalancing eroded the absolute performance, yielding a Net Return of 1396.80% (Annualized Net Sharpe 1.0121).

Conversely, the highly flexible non-linear models significantly underperformed. The **Polynomial Kernel** generated the weakest net results (586.72%) and the highest monthly turnover (69.40%). This phenomenon is deeply documented in institutional quantitative research: over relatively short rolling windows (120 months) on a concentrated universe of 10 assets, highly parameterized non-linear models are highly susceptible to overfitting idiosyncratic market noise. This overfitting translates operationally into erratic, noise-driven forecast adjustments month-over-month. The resulting portfolio instability drives excessive turnover, destroying theoretical alpha through transaction costs. The Linear model's structural rigidity acts as an organic regularizer, preventing it from learning false patterns and allowing it to generalise signals much more efficiently during out-of-sample regime shifts.

5.3. Survivorship Bias and the Structural Dominance of the 1/N Benchmark

Despite the predictive resilience of the Linear Kernel relative to the non-linear algorithms, the passive **Equal-Weight ($1/N$) benchmark** completely dominated the empirical test, delivering an unparalleled 2679.07% Net Cumulative Return and the highest risk-adjusted performance (Annualized Net Sharpe 1.2055).

This profound outperformance is not an anomaly, but the result of three compounding structural factors. First, the $1/N$ strategy benefits entirely from the long-term equity risk premium without incurring the estimation errors inherent in estimating the Mean-Variance covariance matrix. Second, by holding a strictly passive, buy-and-hold allocation, it incurs exactly 0.00% monthly turnover, completely insulating its compounding capital from the execution drag that heavily penalizes the active ML models.

Third, and most crucially, our deliberate but static universe of 10 mega-cap equities introduces a massive *survivorship bias*. By evaluating ten companies in 2024 that have successfully survived and dominated the global economy since 1990 (e.g., Apple, Microsoft, Amazon), the backtest is conducted on an ex-post selected portfolio of historical mega-winners. Over a 21-year secular bull market, a passive allocation to these surviving giants is mathematically predisposed to outperform. Any active Long/Short machine learning model will frequently identify mean-reversion signals and take short positions against these assets, incurring massive friction costs and heavy directional losses when shorting assets that are structurally destined for exponential growth.

5.4. Active Portfolio Allocation and Cross-Sectional Stock Picking

To illustrate the precise operationalization of the machine learning forecasts, Figure 3 details the specific Long/Short asset weights dynamically computed by the Mean-Variance optimizer using the Gaussian RBF predictions for the final month of the simulation (November 2024).




Figure 3: Mean-Variance Portfolio Allocation for November 2024 (Gaussian RBF Model). The gross exposure is structurally normalized to 100%, allowing the algorithm to dynamically scale active Long/Short positions across the universe based on predicted risk-adjusted returns.

Unlike the static, unconditional 10% allocation mandated by the Equal-Weight benchmark, the KRR framework performs highly active stock picking. Driven by the Bayesian-optimized covariance matrix regularization (γ), the model dynamically identifies specific assets to heavily overweight (e.g., allocating substantial long capital to MSFT and AAPL) while severely underweighting or shorting assets predicted to underperform relative to their volatility. While this dispersion confirms that the algorithm successfully extracts active cross-sectional mispricings, the long-term cumulative equity curves demonstrate that high-turnover active shorting in a survivorship-biased universe ultimately fails to overcome the gravitational pull of execution costs and the equity risk premium.

6. ECONOMIC DRIVERS: FEATURE IMPORTANCE

6.1. Methodology: The Information Coefficient (IC)

A frequent and valid criticism of Machine Learning in quantitative finance is its inherent "black box" nature, which obscures the economic intuition behind its algorithmic allocations. To extract transparent economic insights and determine exactly which characteristics drive cross-sectional returns over our 21-year horizon, we evaluate the raw predictive power of each individual factor using the Mean Information Coefficient (IC).

In quantitative asset pricing, the IC is defined as the cross-sectional Spearman rank correlation between the standardized characteristic value at time t ($Z_{i,c,t}$) and the realized asset return at time $t + 1$ ($R_{i,t+1}$). We deliberately employ the Spearman rank correlation rather than the standard Pearson correlation. Financial asset returns are notoriously non-normal and heavily right-skewed; Spearman correlation evaluates monotonic rank relationships rather than strict linear ones, rendering the metric highly robust to extreme market outliers.

Crucially, the IC is an entirely model-agnostic metric. It measures the pure, unconditional predictive alpha of the raw engineered features before any Kernel Ridge Regression algorithm or Mean-Variance portfolio constraint is applied.




Figure 4: Feature Importance: Mean Information Coefficient (IC). Green bars indicate a positive rank correlation (higher factor values predict higher future returns), while red bars indicate a negative correlation (higher factor values predict lower future returns).

6.2. Economic Interpretation of the Predictive Signals

Figure 4 illustrates the empirical Mean IC for the eight engineered characteristics across the entire dataset. The magnitude and direction of these correlations offer profound insights into the market mechanics of our selected large-cap universe.

- **Trading Volume and Volatility (Strongest Positive Drivers):** *Log Volume* (IC = 0.041) and *Historical Volatility* (IC = 0.040) emerge as the most powerful positive predictors of future returns. In this specific universe of mega-cap equities, elevated historical volatility strongly signals higher subsequent returns. This can be interpreted mathematically as a classic risk premium mechanism: investors demand systematically higher expected compensation to hold highly volatile assets. Similarly, the high positive IC for volume suggests a liquidity or attention premium, where highly traded stocks continue to attract capital inflows.
- **Momentum and Trend-Following Anomaly:** *Max Return* (IC = 0.029) and *12-Month Momentum* (IC = 0.025) exhibit robust positive correlations. The positive IC for the 12-month momentum factor empirically confirms the persistent cross-sectional "trend-following" anomaly within equity markets: stocks that have consistently appreciated over the past year tend to continue winning. Interestingly, *Long-Term Reversal* (IC = 0.025) also acts positively in this specific mega-cap subset, suggesting that historical long-term winners in this survivorship-biased universe tend to maintain their trajectory rather than reverting.
- **Microstructural Frictions (Negative Drivers):** Conversely, the *Amihud Illiquidity* proxy (IC = -0.011) and *Short-term Reversal* (IC = -0.008) act as negative drivers. The negative coefficient for the Amihud proxy implies that illiquidity spikes in highly liquid mega-caps often accompany severe negative price shocks, leading to immediate future underperformance. Furthermore, the negative Short-term Reversal formally confirms short-term microstructural mean-reversion: assets that experienced an extreme, unjustified price spike in the immediate prior month are statistically likely to face structural corrections (profit-taking or dealer inventory unwinding) in the subsequent month.

7. CONCLUSION

This study rigorously evaluated the predictive efficacy of Machine Learning in cross-sectional asset pricing, contrasting the out-of-sample performance of individual Kernel Ridge Regression (KRR) architectures against a naive Equal-Weight ($1/N$) benchmark. By processing a deep historical panel dataset of major U.S. equities spanning over three decades, we engineered eight distinct characteristics to forecast forward-looking returns.

The primary contribution of this paper lies in its strict adherence to methodological tractability and institutional-grade realism. We systematically eliminated look-ahead bias through cross-sectional standardization and abandoned static dataset splitting in favor of a rigorous Walk-Forward (Rolling Window) optimization. Furthermore, by implementing Bayesian Optimization via the Tree-structured Parzen Estimator (TPE) algorithm, the models dynamically recalibrated their hyperparameter configurations to maximize friction-adjusted performance (Annualized Net Sharpe Ratio) across vastly different economic regimes.

Our empirical results over a 21-year out-of-sample period (2003-2024) offer profound insights into the mechanics of algorithmic portfolio management. First, we observed a severe bias-variance tradeoff compounded by market frictions. Over 120-month rolling windows on a concentrated asset universe, non-linear algorithms demonstrated a chronic tendency to overfit idiosyncratic market noise. This overfitting generated highly unstable forecasts, translating into excessive portfolio turnover (averaging 69.40% monthly for the Polynomial model) that rapidly destroyed theoretical alpha through transaction costs. Conversely, the structurally rigid Linear Kernel generalized significantly better, yielding the most robust net predictive performance among the active strategies (Net Cumulative Return of 1396.80%).

Secondly, despite the signal extraction capabilities of the linear ML model, the parameter-free Equal-Weight benchmark heavily dominated the empirical test, achieving an unparalleled 2679.07% Net Cumulative Return. This dominance explicitly highlights a crucial limitation frequently overlooked in algorithmic backtests: *survivorship bias*. Implementing a high-turnover, active Long/Short strategy on a static, ex-post selected universe of historical mega-winners structurally detracts from the underlying equity risk premium that a zero-turnover, passive long-only strategy fully captures.

Future research should aim to bridge these gaps by deploying this pipeline on a dynamic, point-in-time asset universe (e.g., a historically accurate rolling S&P 500 constituency) to completely eradicate survivorship bias. Furthermore, incorporating direct turnover penalties (L1 regularization constraints) into the Mean-Variance optimization objective function could further tame the execution drag of active models. Nonetheless, this paper demonstrates that mathematically rigorous Machine Learning frameworks—when appropriately regularized, strictly evaluated out-of-sample net of trading costs, and subjected to transparent economic feature analysis—possess profound analytical value in decoding the noisy cross-section of global financial markets.